

PYTHON PROGRAMMING AND MACHINE LEARNING

INTRODUCTION TO TEXT PROCESSING

Yunghans Irawan (yirawan@nus.edu.sg)

Objectives

- Understand the basic tasks of text processing in Python
- Able to implement a simple text processing using machine learning

Text Processing

- Structured vs. Unstructured Data
- Text Data Preparation
 - Tokenization
 - Stemming / Lemmatization
 - Stop words
- Text Featurization
 - Count Vectorization
 - TF-IDF
 - Word Embeddings

Why Text Processing?

Text data is string with varying lengths

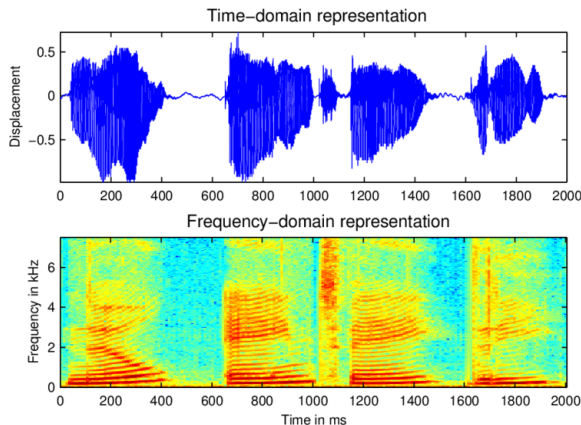
Machine Learning applies **mathematical models**

Therefore, need to **convert text to numbers** via

- Text Processing
- Text Featurization

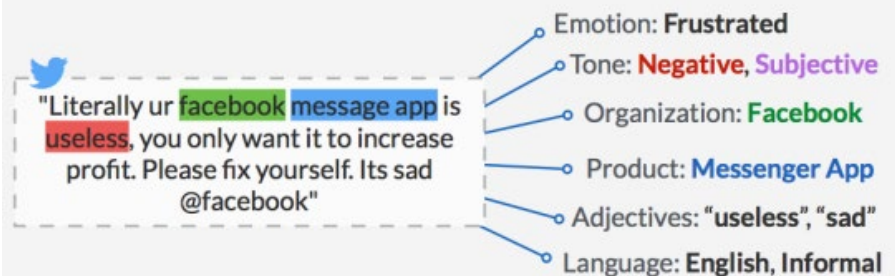
Structured vs. Unstructured Data

| | mpg | cylinders | displacement | horsepower | weight | acceleration | model_year | origin |
|---|------|-----------|--------------|------------|--------|--------------|------------|--------|
| 0 | 18.0 | 8 | 307.0 | 130.0 | 3504.0 | 12.0 | 70 | 1 |
| 1 | 15.0 | 8 | 350.0 | 165.0 | 3693.0 | 11.5 | 70 | 1 |
| 2 | 18.0 | 8 | 318.0 | 150.0 | 3436.0 | 11.0 | 70 | 1 |
| 3 | 16.0 | 8 | 304.0 | 150.0 | 3433.0 | 12.0 | 70 | 1 |
| 4 | 17.0 | 8 | 302.0 | 140.0 | 3449.0 | 10.5 | 70 | 1 |



“zero, one, two, three”

Understanding Language



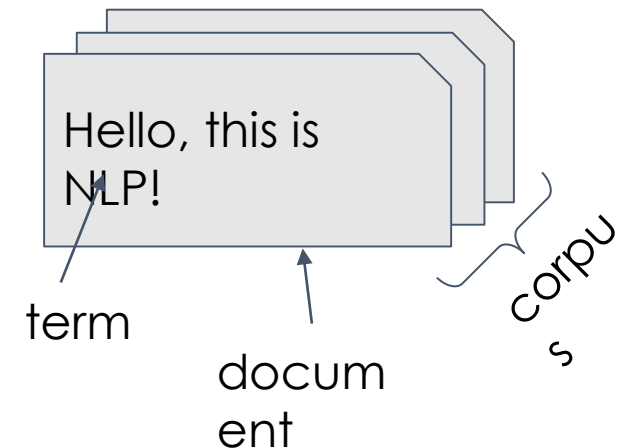
Structured vs. Unstructured Data

| Structured | Semi-structured | Unstructured |
|---------------------|-------------------------|----------------------------|
| Fixed format / size | Data with semantic tags | No format |
| Tabular data | XML, JSON | Text, Audio, Video, Speech |

Objective: Convert unstructured data into structured vectors

Terminology

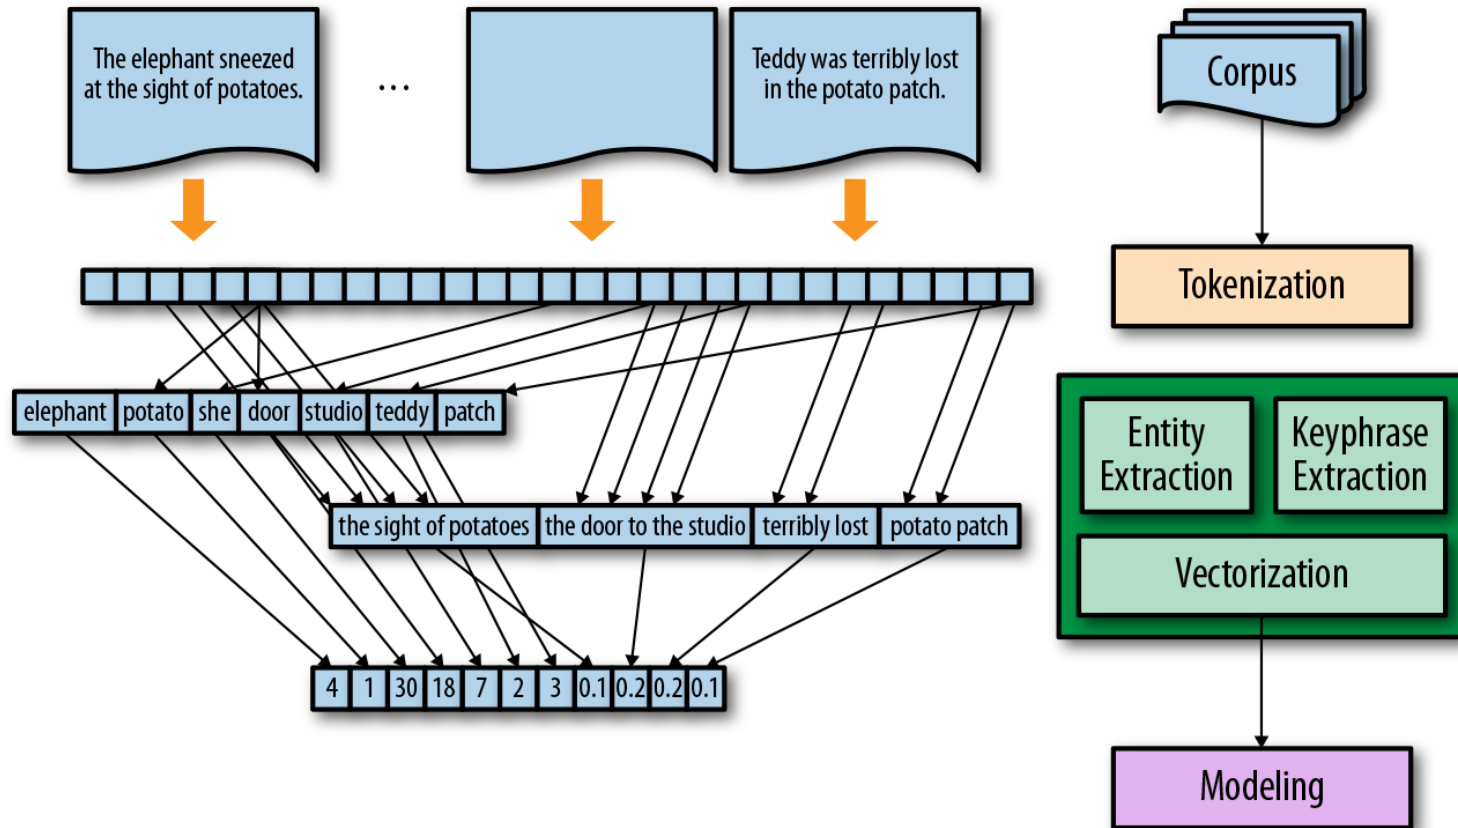
- Document
 - Contains a collection of words / bag of words
 - Article, email, SMS, word document, sentence, ...
- Corpus
 - Collection of documents
- Term / word / token
 - Text entity
- N-gram
 - **Terms** consisting of N consecutive, overlapping sequences of words
- Vocabulary
 - **Set** of unique terms
 - Feature dimensions = vocabulary size



Vocabulary (list of 3 bi-grams):

- hello this
- this is
- is nlp

Text Processing Example



Text Processing Libraries

NLP libraries

- Natural Language Toolkit (NLTK)
- SpaCy
- Gensim

More complete,
linguistic features

ML libraries

- Scikit-Learn
- Keras
- PyTorch

Simple but more
limited

TEXT DATA PREPARATION

Text Data Preparation

- Cleaning
- Tokenization
- Stemming / Lemmatization
- Stop words

Cleaning

- Remove punctuation
- Convert all to lowercase
- Remove non-ASCII characters
- ...

Note: choose the appropriate cleaning for your task

Note: language specific rules

Tokenization

- Split document into terms
- Use libraries or write custom regex

```
from nltk import word_tokenize
```

```
text = 'Hello this is a test.'
```

```
word_tokenize(text)|
```

```
['Hello', 'this', 'is', 'a', 'test', '.']
```

Lemmatization vs Stemming

```
from nltk.stem import WordNetLemmatizer  
  
text = 'he liked cats and dogs, and teaching machines to learn'  
  
lm = WordNetLemmatizer()  
  
print([lm.lemmatize(token) for token in word_tokenize(text)])
```

```
['he', 'liked', 'cat', 'and', 'dog', ',', 'and', 'teaching', 'machine', 'to', 'learn']
```

Looks at word form
(verb, noun, ...)

```
from nltk.stem import SnowballStemmer  
  
text = 'he liked cats and dogs, and teaching machines to learn'  
  
stem = SnowballStemmer(language='english')  
  
print([stem.stem(token) for token in word_tokenize(text)])
```

```
['he', 'like', 'cat', 'and', 'dog', ',', 'and', 'teach', 'machin', 'to', 'learn']
```

Blunt knife chops off
affixes for any word

Usually
either is ok
(depends
on your
task)

Stop words

Stop words are words that are very commonly in use in any sentence

Usually can be **removed without changing meaning**

English stop words

```
{'very', 'itself', 'does', 'nor', 'as', 'had', 'not', 'ours', "shan't", 'out', 'yourself', "hadn't", "hasn't", 'him', 'ma', 'over', 'each', 'is', "that'll", 'she', 'to', "she's", 'but', 'should', 'shouldn', 'needn', 'when', 'those', "weren't", 'don', 'didn', 's', 'if', 'did', 'into', 'more', 'no', 'it', 'doing', "didn't", 'these', 'just', 'then', 'what', 'a', 'ain', 'now', 've', "mightn't", 'his', 'them', 'up', 'he', 'was', 'won', "won't", 'such', 'wasn', 'were', 'theirs', 'or', 'from', 'yours', 'needn't', 'few', 'once', 'd', 'can', 'during', 'they', 'own', 'will', "haven't", "isn't", 'there', 'some', 'y', 'at', 'on', "don't", 'we', "you'd", 'against', 'both', 'aren', 'shan', 're', 'himself', 'be', 'have', 'being', 'hadn', 'any', "wouldn't", 'of', 'under', 'why', 'which', 'after', 'has', 'between', 'again', 'further', 'me', 'do', 'all', 'you', 'and', 'same', 'so', 'than', "you've", 'down', 'weren', 'an', 'most', 'couldn', 'o', 'are', "wasn't", 'who', 'because', 'her', 'before', 'wouldn', 'mightn', 'its', 'this', 'for', "you're", 'i', 'with', 'here', 'above', "should've", "couldn't", 'ourselves', 'where', 'm', 'other', 'in', 'by', 'yourselves', 'themselves', 'hasn', "mustn't", "it's", 'off', "aren't", 'the', 'doesn', 'through', 't', 'you', 'un', "you'll", 'herself', 'whom', 'mustn', 'that', 'am', 'until', 'isn', 'having', 'how', 'about', 'll', 'haven', 'myself', 'hers', 'my', 'while', "doesn't", 'our', 'only', "shouldn't", 'their', 'below', 'too', 'been'}
```

Chinese stop words

Stop words

```
from nltk.corpus import stopwords
stop = set(stopwords.words('english'))

text = 'he liked cats and dogs, and teaching machines to learn'

print([token for token in word_tokenize(text) if token not in stop])

['liked', 'cats', 'dogs', ',', 'teaching', 'machines', 'learn']
```

Note: customize stop words depending on your task

Note: alternative is to apply a maximum threshold on word frequency

TEXT FEATURIZATION

Common Featurization Methods

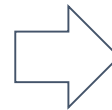
- Word Count
- TF-IDF
- Word Embeddings

Objective: Convert a word to a **meaningful number** or a **vector of numbers**

Word Count Vectorization

1. Create vocabulary from the unique words
2. Count how often each word appears in a document
3. Create a feature vector with the word count as the entry

```
corpus = [  
    'This is the first document.',  
    'This document is the second document.',  
    'And this is the third one.',  
    'Is this the first document?',  
]
```



```
vocabulary = [  
    'and', 'document', 'first', 'is',  
    'one', 'second', 'the', 'third', 'this'  
]
```

Word Count Vectorization

1. Create vocabulary from the unique words
2. Count how often each word appears in a document
3. Create a feature vector with the word count as the entry

| and | document | first | is | one | second | the | third | this | text |
|-----|----------|-------|----|-----|--------|-----|-------|------|---------------------------------------|
| 0 | 1 | 1 | 1 | 0 | 0 | 1 | 0 | 1 | This is the first document. |
| 0 | 2 | 0 | 1 | 0 | 1 | 1 | 0 | 1 | This document is the second document. |
| 1 | 0 | 0 | 1 | 1 | 0 | 1 | 1 | 1 | And this is the third one. |
| 0 | 1 | 1 | 1 | 0 | 0 | 1 | 0 | 1 | Is this the first document? |

Feature vector

Assume more frequent = more important

Address shortcoming of Word Count Vectorization

- Penalize words that occur in lots of documents
- If a word appears all the time, it does not contain much information

Combines two measures

- Term Frequency = Word Count (as before)
- Inverse Document Frequency = Count of documents containing word

TF-IDF (textbook definition)

Word count (number of times
term t appears in document d)



$$\text{tf-idf}(t,d) = \text{tf}(t,d) \times \text{idf}(t)$$

$$\text{idf}(t) = \log \frac{n_d}{1 + \text{df}(d,t)}$$



Document count (number
of documents that contain
term t)

TF-IDF vs. Word Count

Word
Count

| and | document | first | is | one | second | the | third | this | text |
|-----|----------|-------|----|-----|--------|-----|-------|------|---------------------------------------|
| 0 | 1 | 1 | 1 | 0 | 0 | 1 | 0 | 1 | This is the first document. |
| 0 | 2 | 0 | 1 | 0 | 1 | 1 | 0 | 1 | This document is the second document. |
| 1 | 0 | 0 | 1 | 1 | 0 | 1 | 1 | 1 | And this is the third one. |
| 0 | 1 | 1 | 1 | 0 | 0 | 1 | 0 | 1 | Is this the first document? |

TF-IDF

| and | document | first | is | one | second | the | third | this | text |
|----------|----------|----------|----------|----------|----------|----------|----------|----------|---------------------------------------|
| 0.000000 | 0.469791 | 0.580286 | 0.384085 | 0.000000 | 0.000000 | 0.384085 | 0.000000 | 0.384085 | This is the first document. |
| 0.000000 | 0.687624 | 0.000000 | 0.281089 | 0.000000 | 0.538648 | 0.281089 | 0.000000 | 0.281089 | This document is the second document. |
| 0.511849 | 0.000000 | 0.000000 | 0.267104 | 0.511849 | 0.000000 | 0.267104 | 0.511849 | 0.267104 | And this is the third one. |
| 0.000000 | 0.469791 | 0.580286 | 0.384085 | 0.000000 | 0.000000 | 0.384085 | 0.000000 | 0.384085 | Is this the first document? |

More weight to rare words like "first, third"

Word Embeddings

Word count and TF-IDF are **statistical models**

Word Embeddings is a **neural probabilistic model**

Objectives:

1. Infer meaning of a word in terms of its neighbours
2. Compress sparse, high-dimensional data into lower dimensions

Sparsity and high-dimension

Each word is a column

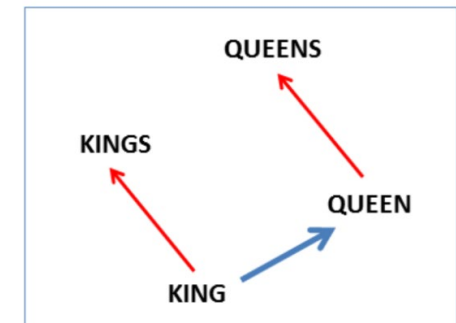
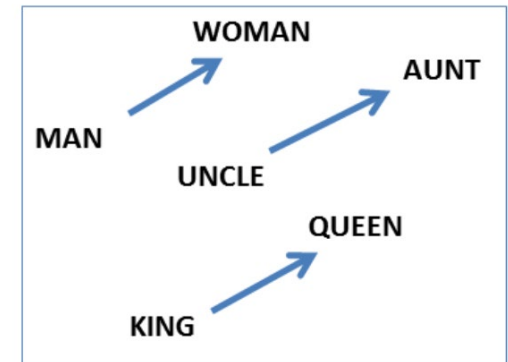
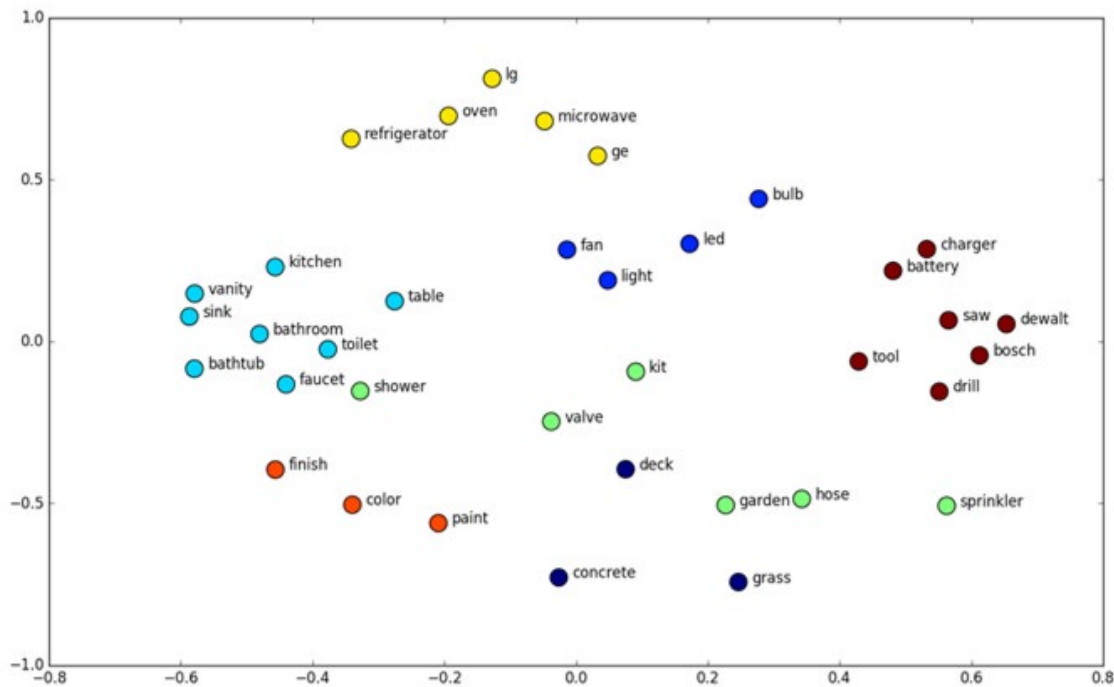
- E.g. 1000 word vocabulary => 1000 columns => 1000 dimensions

Columns are sparse: contain a lot of zeros

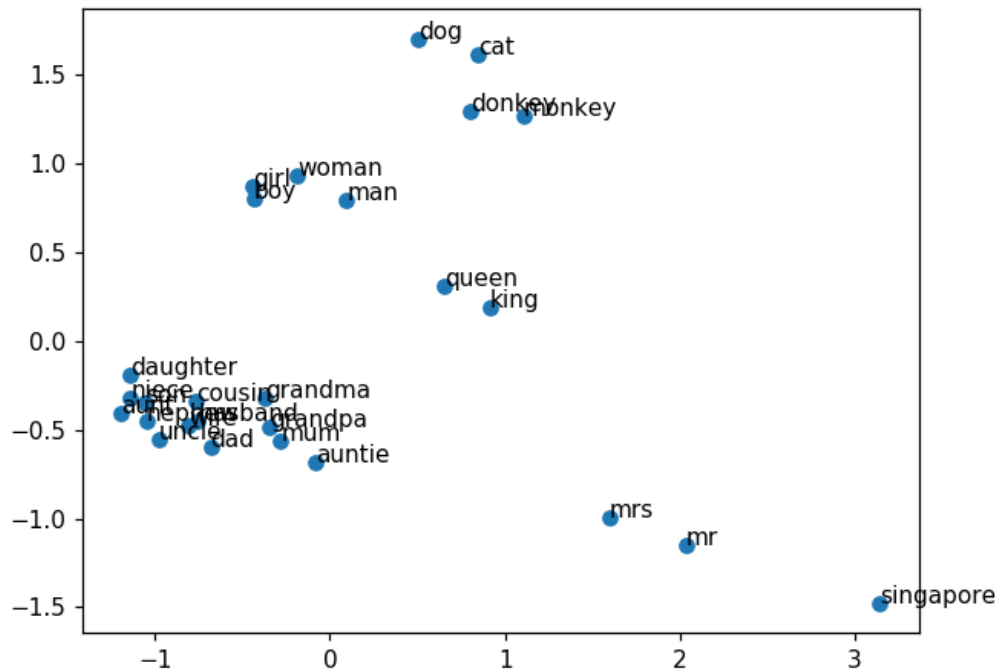
- Most words don't appear frequently

| and | document | first | is | one | second | the | third | this | text |
|----------|----------|----------|----------|----------|----------|----------|----------|----------|---------------------------------------|
| 0.000000 | 0.469791 | 0.580286 | 0.384085 | 0.000000 | 0.000000 | 0.384085 | 0.000000 | 0.384085 | This is the first document. |
| 0.000000 | 0.687624 | 0.000000 | 0.281089 | 0.000000 | 0.538648 | 0.281089 | 0.000000 | 0.281089 | This document is the second document. |
| 0.511849 | 0.000000 | 0.000000 | 0.267104 | 0.511849 | 0.000000 | 0.267104 | 0.511849 | 0.267104 | And this is the third one. |
| 0.000000 | 0.469791 | 0.580286 | 0.384085 | 0.000000 | 0.000000 | 0.384085 | 0.000000 | 0.384085 | Is this the first document? |

Word Meaning in Vector Space



Google News Embeddings



```
model['man']
```

```
array([ 0.32617188,  0.13085938,  0.03466797, -0.08300781,  0.08984375,
        -0.04125977, -0.19824219,  0.00689697,  0.14355469,  0.0019455 ,
         0.02880859, -0.25        , -0.08398438, -0.15136719, -0.10205078,
         0.04077148, -0.09765625,  0.05932617,  0.02978516, -0.10058594,
        -0.13085938,  0.001297   ,  0.02612305, -0.27148438,  0.06396484,
        -0.19140625, -0.078125   ,  0.25976562,  0.375       , -0.04541016,
         0.16210938,  0.13671875, -0.06396484, -0.02062988, -0.09667969,
         0.25390625,  0.24804688, -0.12695312,  0.07177734,  0.3203125 ,
         0.03149414, -0.03857422,  0.21191406, -0.00811768,  0.22265625,
        -0.13476562, -0.07617188,  0.01049805, -0.05175781,  0.03808594,
        -0.13378906,  0.125       ,  0.0559082 , -0.18261719,  0.08154297,
        -0.08447266, -0.07763672, -0.04345703,  0.08105469, -0.01092529,
         0.17480469,  0.30664062, -0.04321289, -0.01416016,  0.09082031])
```

```
model.distance('man', 'woman')
```

```
0.23359877690046482
```

```
model.most_similar("man")
```

```
[('woman', 0.7664012312889099),
 ('boy', 0.6824870109558105),
 ('teenager', 0.6586930751800537),
 ('teenage_girl', 0.6147903203964233),
 ('girl', 0.5921714305877686),
 ('suspected_purse_snatcher', 0.571636438369751),
 ('robber', 0.5585119128227234),
 ('Robbery_suspect', 0.5584409236907959),
 ('teen_ager', 0.5549196004867554),
 ('men', 0.5489763021469116)]
```

Word Distance

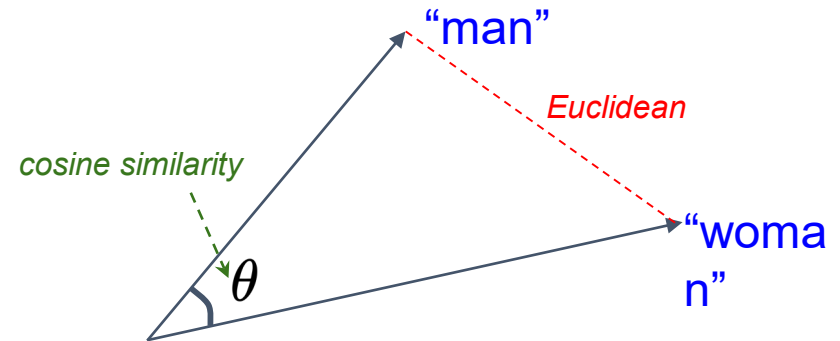
$$\text{similarity}(A, B) = \cos(\theta) = \frac{AB^\top}{\|A\| \|B\|}$$

$$\text{euclidean}(A, B) = \|A - B\| = \sqrt{\|A\|^2 + \|B\|^2 - 2A \cdot B}$$

$$\|A\| = \sqrt{\sum_{i=1}^n A_i^2}$$

More common: cosine similarity

- Has direction
- Range $[-1, 1]$

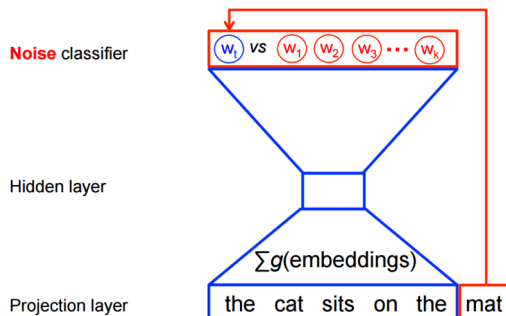


Pre-trained Word Embeddings

- GloVe - trained using global co-occurrence statistics

| Probability and Ratio | $k = solid$ | $k = gas$ | $k = water$ | $k = fashion$ |
|-----------------------|----------------------|----------------------|----------------------|----------------------|
| $P(k ice)$ | 1.9×10^{-4} | 6.6×10^{-5} | 3.0×10^{-3} | 1.7×10^{-5} |
| $P(k steam)$ | 2.2×10^{-5} | 7.8×10^{-4} | 2.2×10^{-3} | 1.8×10^{-5} |
| $P(k ice)/P(k steam)$ | 8.9 | 8.5×10^{-2} | 1.36 | 0.96 |

- Word2Vec - trained using negative sampling



Pre-trained Word Embeddings

- [fastText](#) - each word is a bag of character-level n-grams. Available in 157 languages
- Custom corpus
 - Train your own word-embeddings using [gensim.Word2Vec](#)
 - Alternatively, you can try [updating an existing word embedding](#)

Embedding Layer

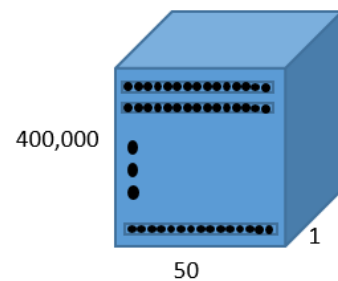
Input Sequence

"I thought the movie was incredible and inspiring"

Integerized Representation

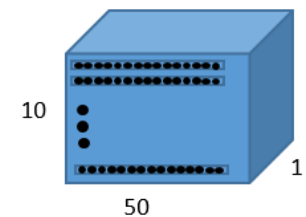
[41 804 201534 1005 15 7446 5 13767 0 0]

Embedding Matrix



`tf.nn.embedding_lookup`

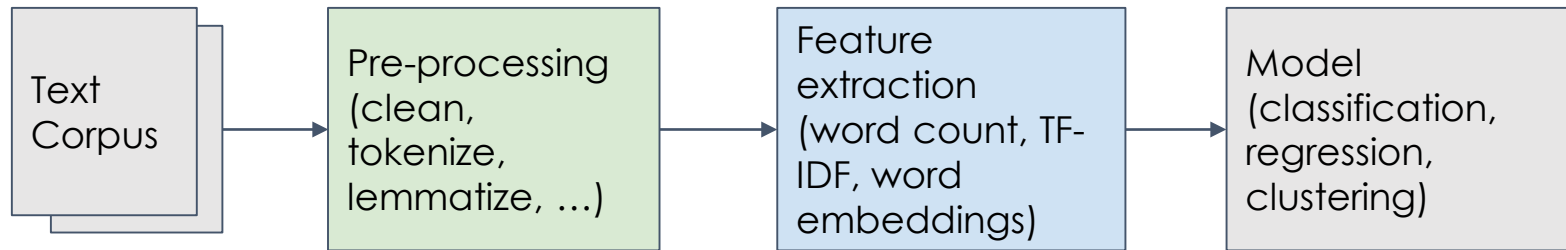
Sequence Vector



When to use...

| | |
|----------------------------|--|
| Word Counts | Small corpus and vocabulary When you need a quick baseline |
| TF-IDF | Documents cover similar vocabulary e.g. movie reviews |
| Pre-trained Word Embedding | Documents cover large, diverse vocabulary e.g. wikipedia articles |
| Custom Word Embedding | Domain-specific vocabulary Want to encode word meanings |

Where are we?



raw
text

terms

vectors

```
from nltk.corpus import stopwords
stop = set(stopwords.words('english'))

text = 'he liked cats and dogs, and teaching machines to learn'

print([token for token in word_tokenize(text) if token not in stop])

['liked', 'cats', 'dogs', ',', 'teaching', 'machines', 'learn']
```

```
model['cat'] # vector
array([-0.96419 , -0.60978 ,  0.67449 ,  0.35113 ,  0.41317 , -0.21241 ,
        1.3796  ,  0.12854 ,  0.31567 ,  0.66325 ,  0.3391  , -0.18934 ,
        -3.325  , -1.1491 , -0.4129 ,  0.2195  ,  0.8706  , -0.50616 ,
        -0.12781 , -0.066965,  0.065761,  0.43927 ,  0.1758  , -0.56058 ,
        0.13529 ], dtype=float32)
```

Applications of Text Processing

Retail: Suggest product pricing based on item description, etc

Finance: Predict stock price movements from market news, etc

Social: Classifying Wikipedia comments for toxicity

Education: Predict acceptance of teacher project proposals

Business: Predict if a Kickstarter project gets funding

Further study

Latent Semantic Analysis - TF-IDF + SVD
dimensionality reduction

King - man + woman is queen, but why?

Gensim Tutorials

Stanford CS224n

How to use Word Embedding Layers in Keras

Practitioner's Guide